



## ----> Importing libraries

```
In [73]: import pandas as pd
import numpy as np

import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')

from sklearn.metrics import accuracy_score,classification_report,confusion_mat
```

## --->Data Extraction

```
In [2]: data = pd.read_csv('C://Datascience Projects/ML/Datasets/Loans_data.csv')
```

## --->Data auditing

```
In [3]: data.shape
```

```
Out[3]: (50000, 10)
```

```
In [4]: data.head()
```

```
Out[4]:   age  income  years_experience  credit_score  loan_amount  employment_type
0    56    128388                  7    658.435992      43858        Salaried
1    69     33280                 13    723.267738      284064  Self-Employed
2    46     22754                 27    741.598799      98582  Self-Employed
3    32     70258                 34    616.773385      282408        Salaried
4    60     63386                  6    756.487805      171925        Salaried
```

```
In [5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   age              50000 non-null   int64  
 1   income            50000 non-null   int64  
 2   years_experience 50000 non-null   int64  
 3   credit_score      50000 non-null   float64 
 4   loan_amount       50000 non-null   int64  
 5   employment_type   50000 non-null   object  
 6   city               50000 non-null   object  
 7   education          50000 non-null   object  
 8   marital_status     50000 non-null   object  
 9   target             50000 non-null   int64  
dtypes: float64(1), int64(5), object(4)
memory usage: 3.8+ MB
```

```
In [6]: data.isna().sum()
```

```
Out[6]: age          0
         income        0
         years_experience 0
         credit_score    0
         loan_amount     0
         employment_type 0
         city            0
         education       0
         marital_status   0
         target           0
         dtype: int64
```

```
In [7]: data.duplicated().sum()
```

```
Out[7]: 0
```

## ---> Exploratory Data Analysis

```
In [8]: data.head()
```

	age	income	years_experience	credit_score	loan_amount	employment_type
0	56	128388		7	658.435992	Salaried
1	69	33280		13	723.267738	Self-Employed
2	46	22754		27	741.598799	Self-Employed
3	32	70258		34	616.773385	Salaried
4	60	63386		6	756.487805	Salaried

```
In [9]: data = data.rename(columns = {"target": "loan_status"})
```

```
In [10]: data.head()
```

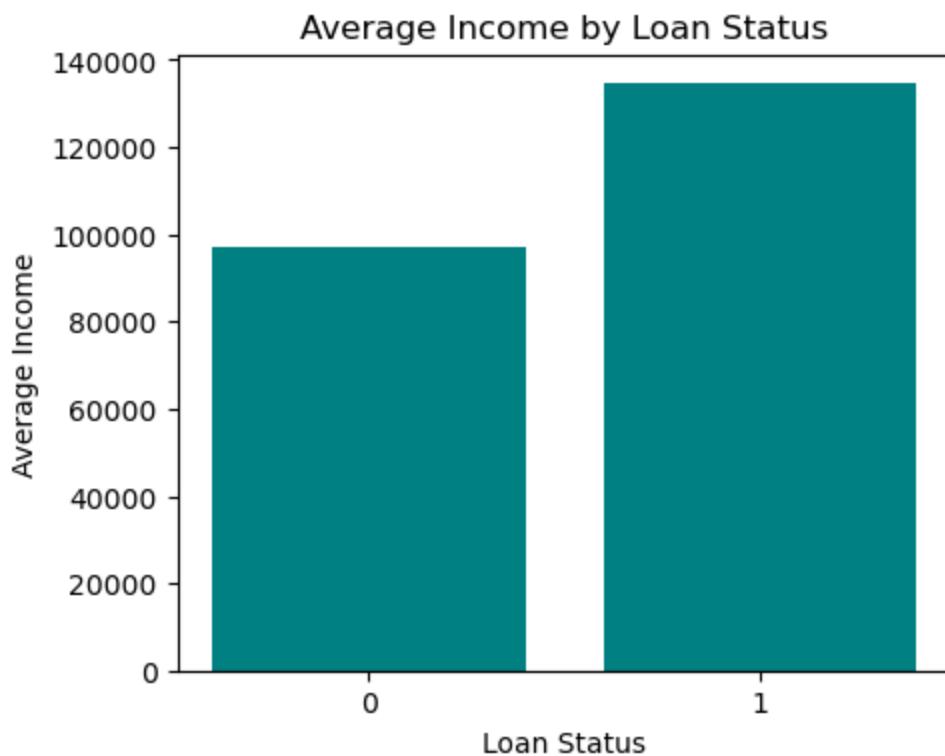
```
Out[10]:   age  income  years_experience  credit_score  loan_amount  employment_type
0    56    128388                  7  658.435992      43858        Salaried
1    69    33280                   13  723.267738      284064  Self-Employed
2    46    22754                   27  741.598799      98582  Self-Employed
3    32    70258                   34  616.773385      282408        Salaried
4    60    63386                   6  756.487805      171925        Salaried
```

```
In [11]: avg_income = data.groupby("loan_status")["income"].mean()
```

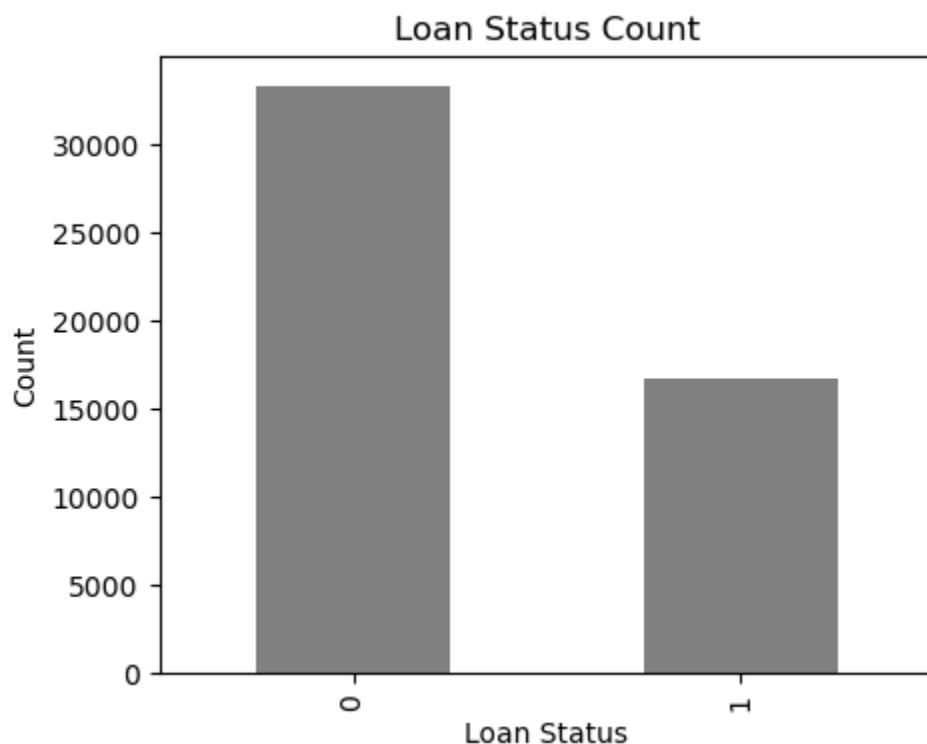
```
In [12]: avg_income
```

```
Out[12]: loan_status
0    97034.294072
1    134526.051963
Name: income, dtype: float64
```

```
In [13]: plt.figure(figsize=(5,4))
plt.bar([0,1], avg_income.values, color = 'teal')
plt.xticks([0, 1], ["0", "1"])
plt.xlabel("Loan Status")
plt.ylabel("Average Income")
plt.title("Average Income by Loan Status")
plt.show()
```



```
In [14]: plt.figure(figsize=(5,4))
data["loan_status"].value_counts().plot(kind="bar", color="grey")
plt.title("Loan Status Count")
plt.xlabel("Loan Status")
plt.ylabel("Count")
plt.show()
```



```
In [15]: data.describe().T
```

		count	mean	std	min
	<b>age</b>	50000.0	43.508260	14.987668	18.000000
	<b>income</b>	50000.0	109545.293680	52043.634207	20002.000000
	<b>years_experience</b>	50000.0	16.956200	10.083428	0.000000
	<b>credit_score</b>	50000.0	649.947228	79.915648	344.399734
	<b>loan_amount</b>	50000.0	252918.827020	142876.555386	5018.000000
	<b>loan_status</b>	50000.0	0.333700	0.471539	0.000000

```
In [16]: data.head()
```

	age	income	years_experience	credit_score	loan_amount	employment_type
<b>0</b>	56	128388		7	658.435992	Salaried
<b>1</b>	69	33280		13	723.267738	Self-Employed
<b>2</b>	46	22754		27	741.598799	Self-Employed
<b>3</b>	32	70258		34	616.773385	Salaried
<b>4</b>	60	63386		6	756.487805	Salaried

```
In [17]: data = data.drop(['city', 'education', 'years_experience', 'marital_status'], axis=1)
```

```
In [18]: data.head()
```

	age	income	credit_score	loan_amount	employment_type	loan_status
<b>0</b>	56	128388	658.435992	43858	Salaried	0
<b>1</b>	69	33280	723.267738	284064	Self-Employed	0
<b>2</b>	46	22754	741.598799	98582	Self-Employed	1
<b>3</b>	32	70258	616.773385	282408	Salaried	0
<b>4</b>	60	63386	756.487805	171925	Salaried	0

```
In [19]: y = data['loan_status']
```

```
In [20]: y.isna().sum()
```

```
Out[20]: 0
```

```
In [21]: num_df = data.select_dtypes(exclude = 'object')
```

```
cat_df = data.select_dtypes(include = 'object')
```

```
In [22]: num_df = num_df.drop(['loan_status'],axis=1)
```

```
In [23]: num_df.head()
```

```
Out[23]:   age  income  credit_score  loan_amount
```

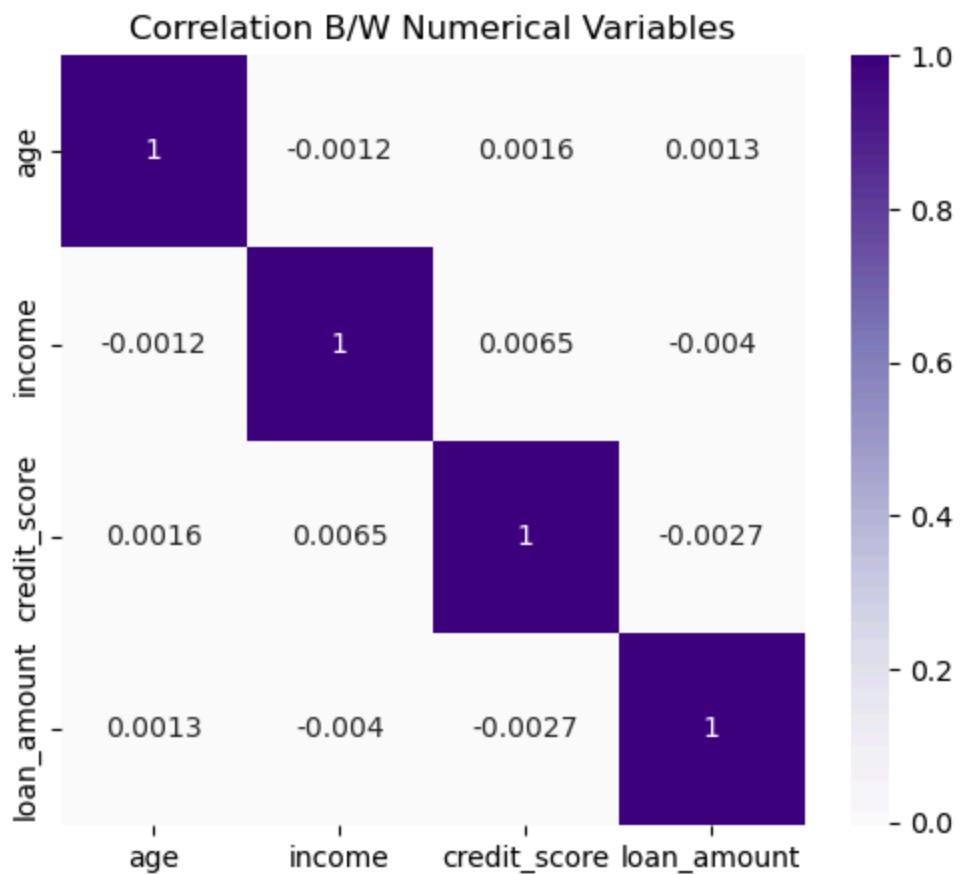
<b>0</b>	56	128388	658.435992	43858
<b>1</b>	69	33280	723.267738	284064
<b>2</b>	46	22754	741.598799	98582
<b>3</b>	32	70258	616.773385	282408
<b>4</b>	60	63386	756.487805	171925

```
In [24]: num_df.corr().T
```

```
Out[24]:   age  income  credit_score  loan_amount
```

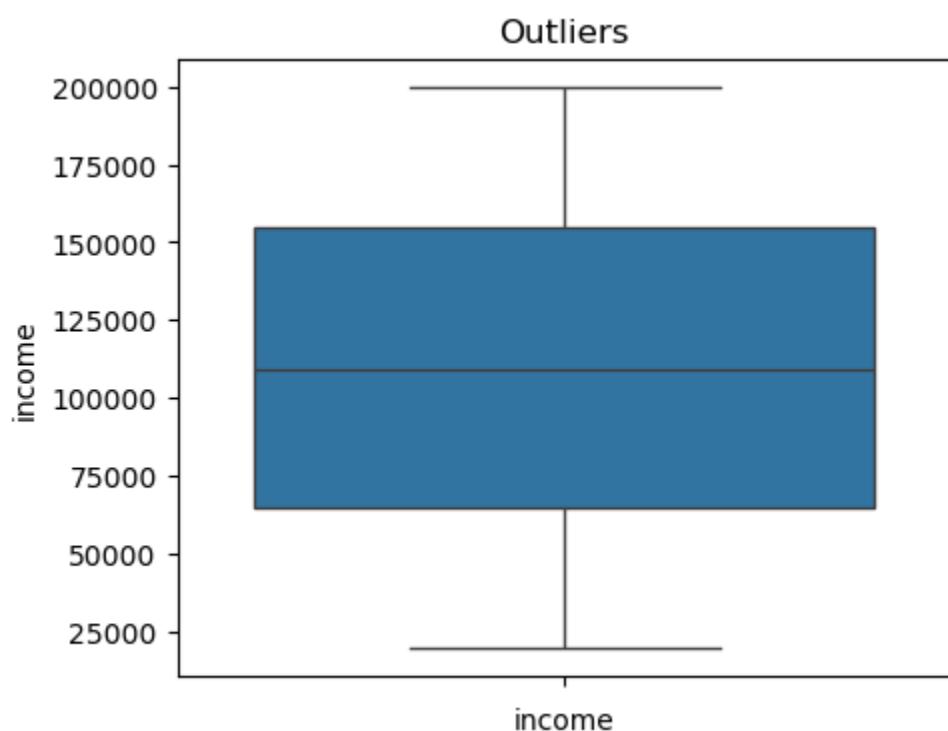
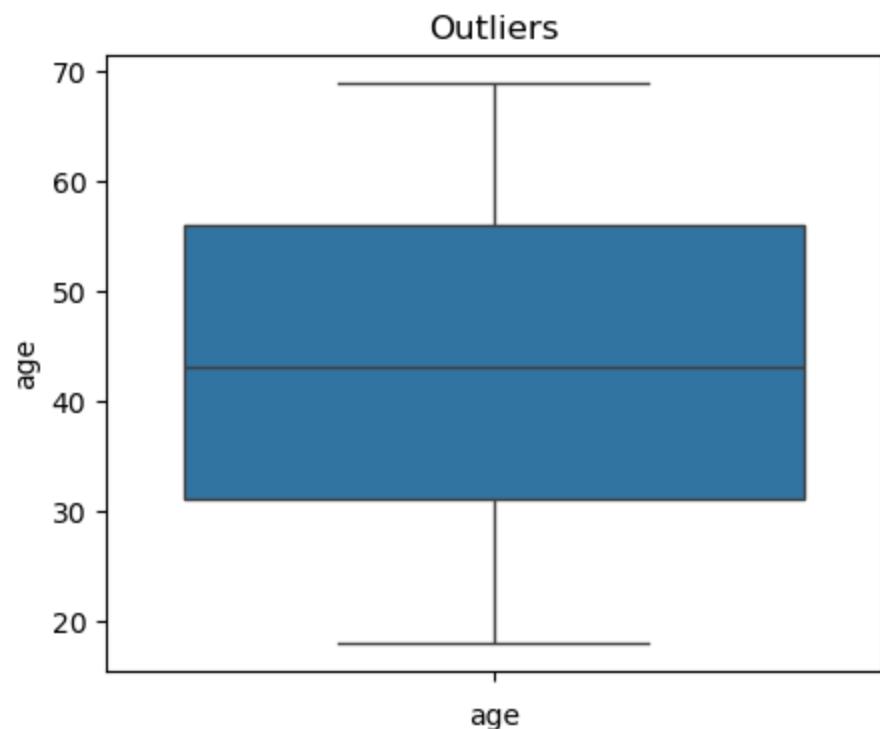
<b>age</b>	1.000000	-0.001247	0.001625	0.001323
<b>income</b>	-0.001247	1.000000	0.006499	-0.004037
<b>credit_score</b>	0.001625	0.006499	1.000000	-0.002746
<b>loan_amount</b>	0.001323	-0.004037	-0.002746	1.000000

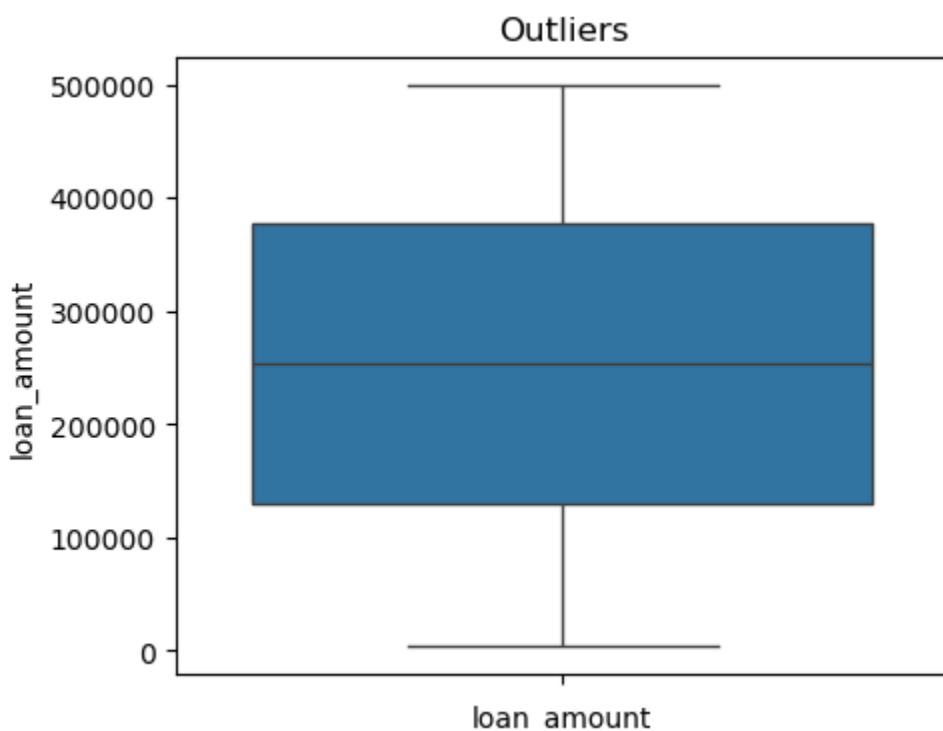
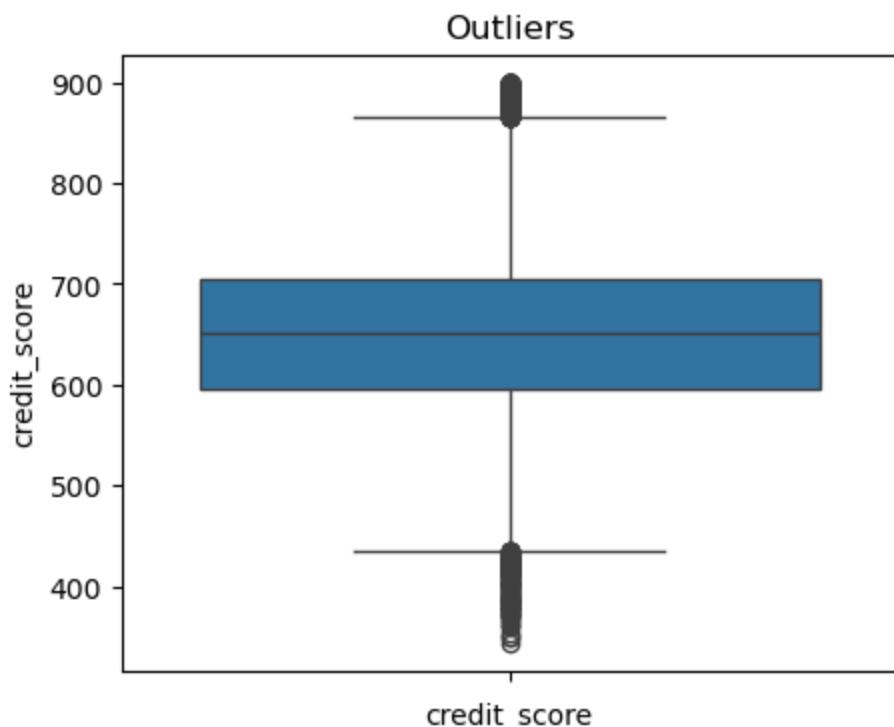
```
In [25]: plt.figure(figsize=(6,5))
sns.heatmap(num_df.corr(), annot=True, cmap='Purples')
plt.title("Correlation B/W Numerical Variables")
plt.show()
```



---> Outlier Checking

```
In [26]: for col in num_df.columns:  
    plt.figure(figsize=(5,4))  
    sns.boxplot(num_df[col])  
    plt.xlabel(col)  
    plt.title("Outliers")  
    plt.show()
```





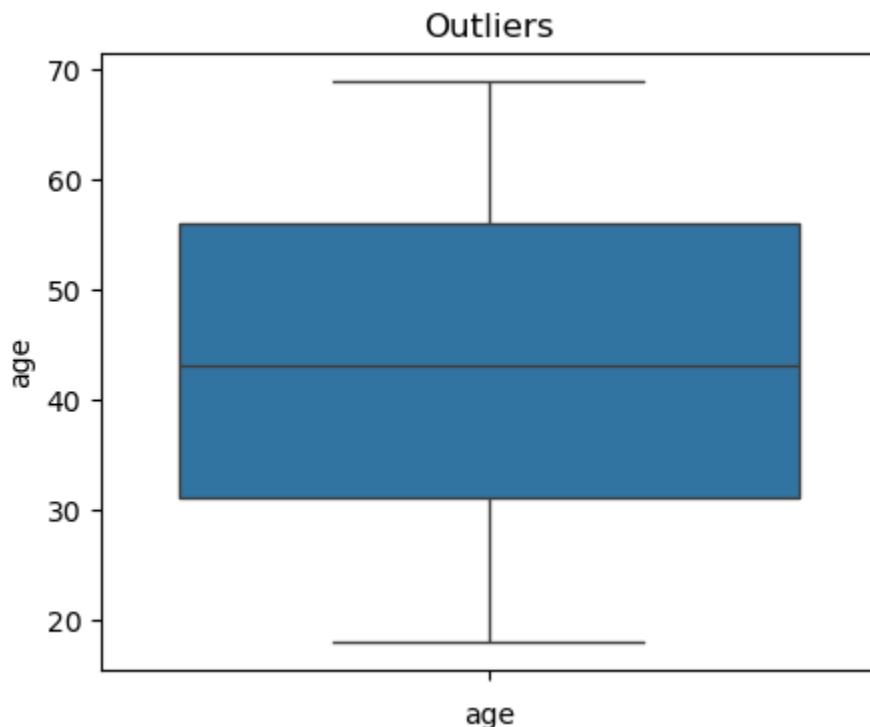
---> Outlier Treating

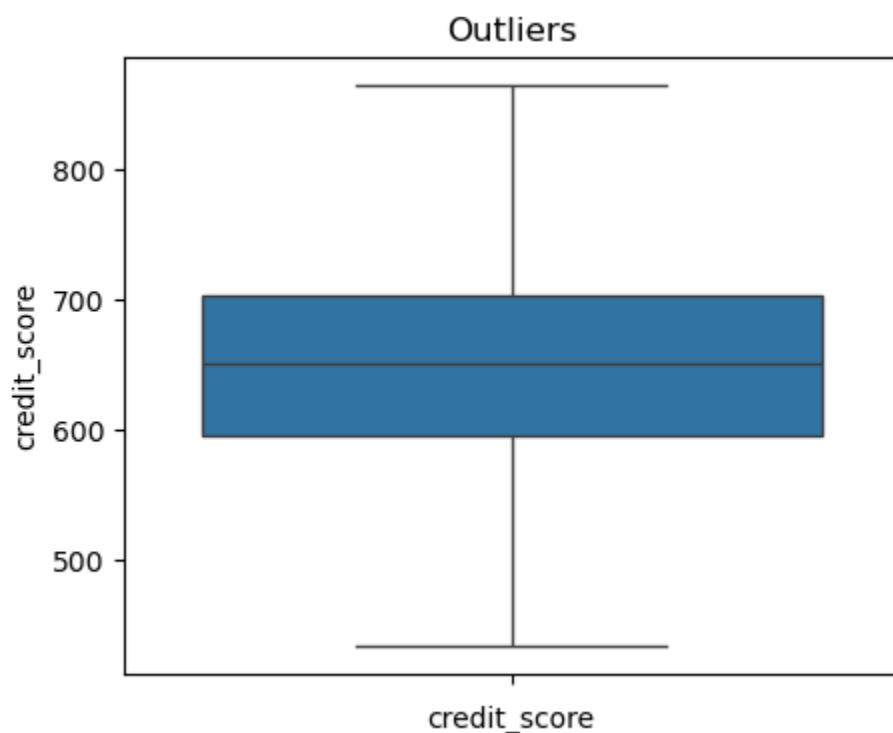
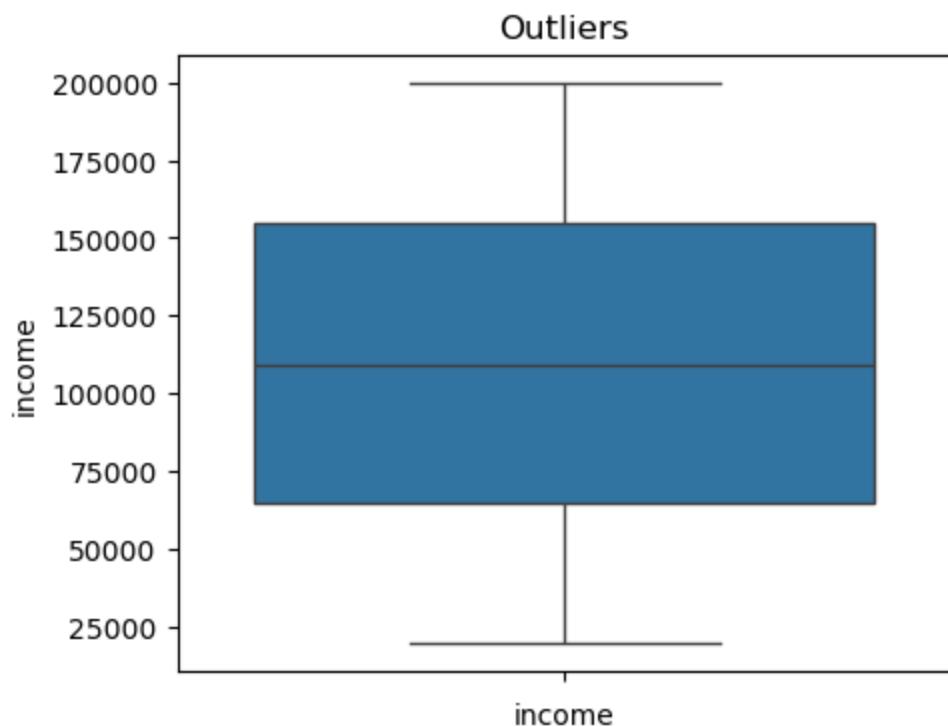
```
In [27]: def outlier_treat(x):
    q1 = x.quantile(0.25)
    q3 = x.quantile(0.75)
    iqr = q3-q1
    uw = q3+1.5*iqr
```

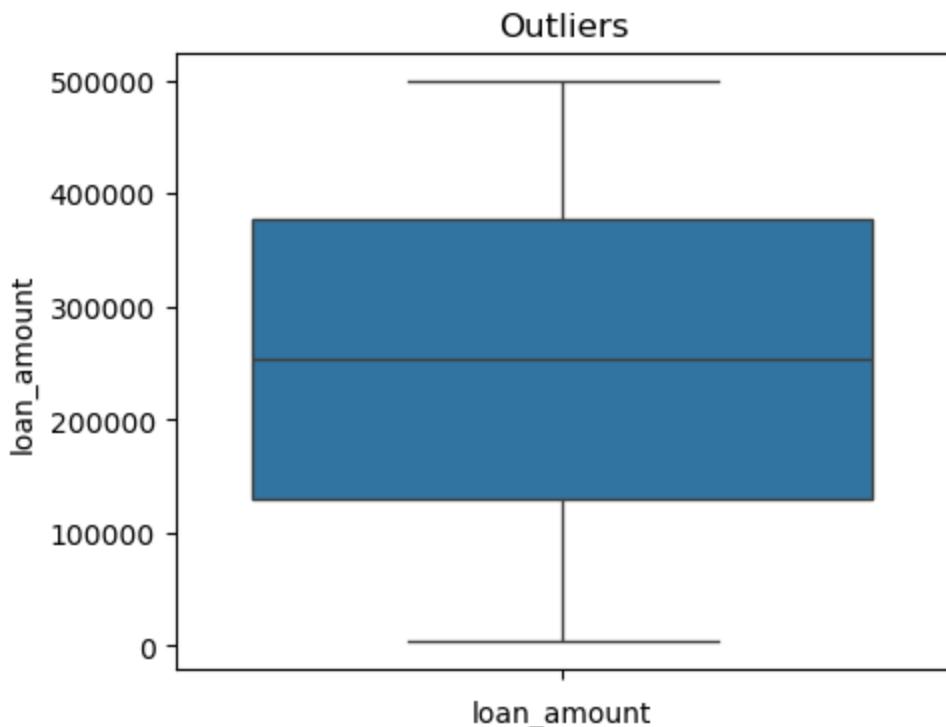
```
    lw = q1-1.5*iqr  
    return(x.clip(upper = uw, lower = lw))
```

```
In [28]: num_df = num_df.apply(outlier_treat)
```

```
In [29]: for col in num_df.columns:  
    plt.figure(figsize=(5,4))  
    sns.boxplot(num_df[col])  
    plt.xlabel(col)  
    plt.title("Outliers")  
    plt.show()
```







### ---> Pre Processing

```
In [30]: from sklearn.preprocessing import LabelEncoder
```

```
In [31]: le = LabelEncoder()  
cat_df["employment_type"] = le.fit_transform(cat_df["employment_type"])
```

```
In [32]: cat_df.head()
```

```
Out[32]: employment_type
```

0	0
1	1
2	1
3	0
4	0

```
In [33]: cat_df['employment_type'].value_counts()
```

```
Out[33]: employment_type  
2    12545  
0    12535  
1    12460  
3    12460  
Name: count, dtype: int64
```

0 → Student

1 → Self-Employed

2 → Salaried

3 → Unemployed

```
In [34]: df = pd.concat([num_df,cat_df],axis=1)
```

```
In [35]: df.head()
```

```
Out[35]:   age  income  credit_score  loan_amount  employment_type
0      56    128388     658.435992       43858          0
1      69     33280     723.267738       284064          1
2      46     22754     741.598799       98582          1
3      32     70258     616.773385       282408          0
4      60     63386     756.487805       171925          0
```

## ---> Data Splitting

```
In [36]: X = df[['age','income','credit_score','loan_amount','employment_type']]
```

```
In [37]: X.head()
```

```
Out[37]:   age  income  credit_score  loan_amount  employment_type
0      56    128388     658.435992       43858          0
1      69     33280     723.267738       284064          1
2      46     22754     741.598799       98582          1
3      32     70258     616.773385       282408          0
4      60     63386     756.487805       171925          0
```

```
In [38]: from sklearn.model_selection import train_test_split
```

```
In [39]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [40]: print(X_train.value_counts())
print(X_test.value_counts())
print(y_train.value_counts())
print(y_test.value_counts())
```

```
age  income  credit_score  loan_amount  employment_type
18   20067    674.656879    248088      1                      1
52   118037   651.707385    144755      3                      1
      116501   596.399498    263253      0                      1
      116552   544.206175    37147       2                      1
      117202   681.804282    370532      3                      1
                                ..
35   46543    671.251005    200503      0                      1
      46393    552.898044    75309       1                      1
      46335    668.462596    452183      1                      1
      45908    612.436593    77777       1                      1
69   199833   627.048528    91015       2                      1
Name: count, Length: 35000, dtype: int64
age  income  credit_score  loan_amount  employment_type
18   20667    675.929909    253120      2                      1
52   196684   668.788866    58391       0                      1
      184069   826.338051    126415      2                      1
      185765   474.724262    90597       1                      1
      187021   703.327453    446742      0                      1
                                ..
35   130532   638.750137    369392      2                      1
      131318   536.860174    372916      3                      1
      131641   660.092813    41481       3                      1
      133795   534.397632    91425       3                      1
69   199821   710.769631    403616      0                      1
Name: count, Length: 15000, dtype: int64
loan_status
0    23323
1    11677
Name: count, dtype: int64
loan_status
0    9992
1    5008
Name: count, dtype: int64
```

## Model Building

### Decision Trees

```
In [41]: from sklearn.tree import DecisionTreeClassifier
```

```
In [42]: dt = DecisionTreeClassifier()
```

```
In [43]: dt.fit(X_train,y_train)
```

```
Out[43]: ▾ DecisionTreeClassifier ⓘ ?  
DecisionTreeClassifier()
```

```
In [47]: dt.predict(X_train)
```

```
Out[47]: array([1, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [48]: accuracy_score(y_train,dt.predict(X_train))
```

```
Out[48]: 1.0
```

```
In [49]: dt.predict(X_test)
```

```
Out[49]: array([1, 0, 0, ..., 1, 0, 0], dtype=int64)
```

```
In [50]: accuracy_score(y_test,dt.predict(X_test))
```

```
Out[50]: 0.8666666666666667
```

## Hyper Tuning

```
In [51]: dt_model = DecisionTreeClassifier(  
        max_depth=6,  
        min_samples_split=10,  
        min_samples_leaf=5,  
        criterion="gini",  
        random_state=42  
)
```

```
In [52]: dt_model.fit(X_train,y_train)
```

```
Out[52]: ▾ DecisionTreeClassifier  
DecisionTreeClassifier(max_depth=6, min_samples_leaf=5, min_samples_s  
plit=10,  
random_state=42)
```

```
In [53]: dt_model.predict(X_train)
```

```
Out[53]: array([1, 0, 0, ..., 1, 0, 0], dtype=int64)
```

```
In [54]: accuracy_score(y_train,dt_model.predict(X_train))
```

```
Out[54]: 0.9157142857142857
```

```
In [55]: print(classification_report(y_train,dt_model.predict(X_train)))
```

```

      precision    recall  f1-score   support

          0       1.00      0.87      0.93     23323
          1       0.80      1.00      0.89     11677

   accuracy                           0.92     35000
macro avg       0.90      0.94      0.91     35000
weighted avg    0.93      0.92      0.92     35000

```

In [56]: `dt_model.predict(X_test)`

Out[56]: `array([1, 0, 0, ..., 1, 1, 0], dtype=int64)`

In [57]: `accuracy_score(y_test,dt_model.predict(X_test))`

Out[57]: `0.9106`

In [58]: `print(classification_report(y_test,dt_model.predict(X_test)))`

```

      precision    recall  f1-score   support

          0       1.00      0.87      0.93     9992
          1       0.79      0.99      0.88     5008

   accuracy                           0.91     15000
macro avg       0.89      0.93      0.90     15000
weighted avg    0.93      0.91      0.91     15000

```

## Ensemble Models

### Gradient Boosting

In [59]: `from sklearn.ensemble import GradientBoostingClassifier  
gbc = GradientBoostingClassifier()  
gbc.fit(X_train,y_train)`

Out[59]: `GradientBoostingClassifier i ?  
GradientBoostingClassifier()`

In [60]: `gbc.predict(X_train)`

Out[60]: `array([1, 0, 0, ..., 1, 0, 0], dtype=int64)`

In [61]: `accuracy_score(y_train,gbc.predict(X_train))`

Out[61]: `0.9158285714285714`

```
In [62]: gbc.predict(X_test)
```

```
Out[62]: array([1, 0, 0, ..., 1, 1, 0], dtype=int64)
```

```
In [63]: accuracy_score(y_test,gbc.predict(X_test))
```

```
Out[63]: 0.9119333333333334
```

## Hyper tuning

```
In [68]: gb_model=GradientBoostingClassifier(learning_rate=0.05,  
n_estimators=300,  
max_depth=3,  
min_samples_split=5,  
min_samples_leaf=2,  
random_state=42)
```

```
In [69]: gb_model.fit(X_train,y_train)
```

```
Out[69]: ▾ GradientBoostingClassifier  
GradientBoostingClassifier(learning_rate=0.05, min_samples_leaf=2,  
min_samples_split=5, n_estimators=300,  
random_state=42)
```

```
In [70]: gb_model.predict(X_train)
```

```
Out[70]: array([1, 0, 0, ..., 1, 0, 0], dtype=int64)
```

```
In [71]: accuracy_score(y_train,gb_model.predict(X_train))
```

```
Out[71]: 0.9159428571428572
```

```
In [72]: print(classification_report(y_train,gb_model.predict(X_train)))
```

	precision	recall	f1-score	support
0	1.00	0.87	0.93	23323
1	0.80	1.00	0.89	11677
accuracy			0.92	35000
macro avg	0.90	0.94	0.91	35000
weighted avg	0.93	0.92	0.92	35000

```
In [74]: print(confusion_matrix(y_train,gb_model.predict(X_train)))
```

```
[[20385  2938]  
 [    4 11673]]
```

```
In [75]: gb_model.predict(X_test)
```

```
Out[75]: array([1, 0, 0, ..., 1, 1, 0], dtype=int64)
```

```
In [76]: print(accuracy_score(y_test,gb_model.predict(X_test)))
```

```
0.9116
```

```
In [77]: print(classification_report(y_test,gb_model.predict(X_test)))
```

	precision	recall	f1-score	support
0	1.00	0.87	0.93	9992
1	0.79	1.00	0.88	5008
accuracy			0.91	15000
macro avg	0.90	0.93	0.91	15000
weighted avg	0.93	0.91	0.91	15000

```
In [78]: print(confusion_matrix(y_test,gb_model.predict(X_test)))
```

```
[[8673 1319]
 [ 7 5001]]
```

## Random Forest

```
In [79]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
```

```
In [80]: rf.fit(X_train,y_train)
```

```
Out[80]: RandomForestClassifier(i ?)
          RandomForestClassifier()
```

```
In [81]: rf.predict(X_train)
```

```
Out[81]: array([1, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [82]: accuracy_score(y_train,rf.predict(X_train))
```

```
Out[82]: 1.0
```

```
In [83]: rf.predict(X_test)
```

```
Out[83]: array([1, 0, 0, ..., 1, 0, 0], dtype=int64)
```

```
In [84]: accuracy_score(y_test,rf.predict(X_test))
```

```
Out[84]: 0.9078
```

## Hyper Tuning

```
In [85]: rf_model = RandomForestClassifier(  
    n_estimators=300,  
    max_depth=12,  
    max_features='sqrt',  
    min_samples_leaf=3,  
    min_samples_split=5,  
    random_state=42  
)
```

```
In [86]: rf_model.fit(X_train,y_train)
```

```
Out[86]: RandomForestClassifier(max_depth=12, min_samples_leaf=3, min_samples_split=5,  
n_estimators=300, random_state=42)
```

```
In [87]: rf_model.predict(X_train)
```

```
Out[87]: array([1, 0, 0, ..., 1, 0, 0], dtype=int64)
```

```
In [88]: accuracy_score(y_train,rf_model.predict(X_train))
```

```
Out[88]: 0.9168857142857143
```

```
In [89]: print(classification_report(y_train,rf_model.predict(X_train)))
```

	precision	recall	f1-score	support
0	1.00	0.88	0.93	23323
1	0.80	1.00	0.89	11677
accuracy			0.92	35000
macro avg	0.90	0.94	0.91	35000
weighted avg	0.93	0.92	0.92	35000

```
In [90]: print(confusion_matrix(y_train,rf_model.predict(X_train)))
```

```
[[20415 2908]  
 [ 1 11676]]
```

```
In [91]: rf_model.predict(X_test)
```

```
Out[91]: array([1, 0, 0, ..., 1, 1, 0], dtype=int64)
```

```
In [92]: accuracy_score(y_test,rf_model.predict(X_test))
```

```
Out[92]: 0.9115333333333333
```

```
In [93]: print(classification_report(y_test,rf_model.predict(X_test)))
```

	precision	recall	f1-score	support
0	1.00	0.87	0.93	9992
1	0.79	1.00	0.88	5008
accuracy			0.91	15000
macro avg	0.90	0.93	0.91	15000
weighted avg	0.93	0.91	0.91	15000

```
In [94]: print(confusion_matrix(y_test,rf_model.predict(X_test)))
```

```
[[8672 1320]
 [ 7 5001]]
```

## XG Boost

```
In [97]: from xgboost import XGBClassifier
xgb = XGBClassifier()
```

```
In [98]: xgb.fit(X_train,y_train)
```

```
Out[98]: XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_roun
              ds=None,
              enable_categorical=False, eval_metric=None, feature_typ
              es=None,
              feature_weights=None, gamma=None, grow_policy=None,
              importance_type=None, interaction_constraints=None,
              learning_rate=None, max_bin=None, max_cat_threshold=None
              e,
```

```
In [99]: xgb.predict(X_train)
```

```
Out[99]: array([1, 0, 0, ..., 1, 0, 0])
```

```
In [100... accuracy_score(y_train,xgb.predict(X_train))
```

```
Out[100... 0.9377142857142857
```

```
In [101... xgb.predict(X_test)
```

```
Out[101... array([1, 0, 0, ..., 1, 1, 0])
```

```
In [102... accuracy_score(y_test,xgb.predict(X_test))
```

```
Out[102... 0.9046
```

## Hyper Tuning

```
In [103... xgb_model = XGBClassifier(  
        n_estimators=300,  
        learning_rate=0.05,  
        max_depth=6,  
        subsample=0.8,  
        colsample_bytree=0.8,  
        gamma=0,  
        reg_lambda=1,  
        random_state=42,  
        n_jobs=-1  
)
```

```
In [104... xgb_model.fit(X_train,y_train)
```

```
Out[104... ▾
```

XGBClassifier

```
XGBClassifier(base_score=None, booster=None, callbacks=None,  
             colsample_bylevel=None, colsample_bynode=None,  
             colsample_bytree=0.8, device=None, early_stopping_round  
s=None,  
             enable_categorical=False, eval_metric=None, feature_typ  
es=None,  
             feature_weights=None, gamma=0, grow_policy=None,  
             importance_type=None, interaction_constraints=None,  
             learning_rate=0.05, max_bin=None, max_cat_threshold=None,  
             max_delta_step=None, max_leaves=None, max_depth=6,  
             min_child_weight=1, missing='NaN', n_estimators=300,  
             n_jobs=-1, nthread=None, objective='binary:  
regression', random_state=42, scale_pos_weight=None,  
             tree_method='auto')
```

```
In [105... xgb_model.predict(X_train)
```

```
Out[105... array([1, 0, 0, ..., 1, 0, 0])
```

```
In [106... accuracy_score(y_train,xgb_model.predict(X_train))
```

```
Out[106... 0.9194857142857142
```

```
In [107... print(classification_report(y_train,xgb_model.predict(X_train)))
```

```

precision    recall   f1-score   support
0            1.00     0.88      0.94      23323
1            0.81     1.00      0.89      11677

accuracy                           0.92      35000
macro avg                           0.90      35000
weighted avg                          0.93      35000

```

In [108]: `print(confusion_matrix(y_train,xgb_model.predict(X_train)))`

```

[[20529 2794]
 [ 24 11653]]

```

In [109]: `xgb_model.predict(X_test)`

Out[109]: `array([1, 0, 0, ..., 1, 1, 0])`

In [110]: `accuracy_score(y_test,xgb_model.predict(X_test))`

Out[110]: `0.9099333333333334`

In [111]: `print(classification_report(y_test,xgb_model.predict(X_test)))`

```

precision    recall   f1-score   support
0            1.00     0.87      0.93      9992
1            0.79     0.99      0.88      5008

accuracy                           0.91      15000
macro avg                           0.89      15000
weighted avg                          0.93      15000

```

In [112]: `print(confusion_matrix(y_test,xgb_model.predict(X_test)))`

```

[[8681 1311]
 [ 40 4968]]

```

## Model Accuracies Table

In [120]: `dt_train_acc = accuracy_score(y_train,dt_model.predict(X_train))`

In [121]: `dt_test_acc = accuracy_score(y_test,dt_model.predict(X_test))`

In [122]: `rf_train_acc = accuracy_score(y_train,rf_model.predict(X_train))`

In [123]: `rf_test_acc = accuracy_score(y_test,rf_model.predict(X_test))`

In [124]: `xgb_train_acc = accuracy_score(y_train,xgb_model.predict(X_train))`

```
In [125... gb_train_acc = accuracy_score(y_train,gb_model.predict(X_train))  
In [126... gb_test_acc = accuracy_score(y_test,gb_model.predict(X_test))  
In [127... xgb_test_acc = accuracy_score(y_test,xgb_model.predict(X_test))  
In [129... results = {  
    "Model": ["Decision Tree", "Random Forest", "XGBoost", "Gradient Boosting"]  
    "Train Accuracy": [dt_train_acc, rf_train_acc, xgb_train_acc, gb_train_acc]  
    "Test Accuracy": [dt_test_acc, rf_test_acc, xgb_test_acc, gb_test_acc]  
}  
  
accuracy_df = pd.DataFrame(results)  
accuracy_df
```

Out[129...]

	Model	Train Accuracy	Test Accuracy
<b>0</b>	Decision Tree	0.915714	0.910600
<b>1</b>	Random Forest	0.916886	0.911533
<b>2</b>	XGBoost	0.919486	0.909933
<b>3</b>	Gradient Boosting	0.915943	0.911600

In [ ]: